

Examining the Relationship of Expenditures and Student Outcomes at Public 4-Year Institutions

Duncan Clark and Leonard Wainstein

March 23, 2018

Introduction

Higher education institutions serve a public good; they educate their students with the hope that in the future these graduates will go on and contribute to society as a whole. However, these students also benefit, as they tend to be more successful than their non university educated peers. Thus, since much of higher education is funded by the public purse, it becomes a matter for the public to decide how much to subsidise individuals' educations for the good of the whole, and in particular it becomes crucial that the public money is used efficiently. It is this efficiency that we concerned ourselves with in this project.

In this project we hoped to analyse quantitatively the relationship between student outcomes at public 4-year higher education institutions and their expenditures. Specifically we looked at 6-year graduation rates for first-time full-time bachelor's degree-seeking student cohorts and the number of bachelor's degrees awarded per full time equivalent student (FTE). We decided to look at two outcomes because certain institutions may be more concerned with its bachelor's degrees per FTE than its graduation rate (or vice versa). A small school with mostly full-time student cohorts would likely be more interested in increasing its graduation rate. However, large schools with many part-time students would also want to know if the rate at which they are producing graduates is keeping up with the total number of credits students are taking that institution.

We were interested in using methods developed in this course to predict the two outcomes, but we were also interested in making valid interpretations of our results. With this in mind we first tried only including variables that we believed to be important as in Tandberg and Hillman (2014) and Research for Action's 2017 study Slaughter et al. (2017). We then tried carrying out automatic variables selection through Lasso regression and using these variables as the basis for our models.

The data we analyzed was obtained from the Delta Cost Project at American Institutes for Research, which gave a dataset for all higher education institutions in the US for 2003 to 2015. The data is publically available at: <https://www.deltacostproject.org/>.

Data Processing

The Delta Cost Project is a panel data set that has institutional characteristics and outcomes for every university in United States and its territories for academic years 1987 to 2015 (academic years are indexed by the year during which conclude - e.g. the 2015 academic year is the 2014-15 school year).

We restricted our sample to 4-year public primarily bachelor's degree-granting institutions in the 50 states (plus DC), and excluded any institutions that did not exist in all years from 2003 to 2015. Years from 2003 onwards generally had a low level of missingness in the outcomes of interest, as well as covariates regarding expenditures, so we chose this as our cutoff for inclusion. We also dropped schools that were missing values in one of the outcomes for any of the years from 2003 to 2015, and schools with missing values for a few other key variables on all years of data. As for examining other types of institutions, this could be interesting

though there is likely to be heterogeneity since the institutions serve different purposes, so for our purposes we simply excluded them. Starting with 553 schools in 2003 to 2015, thinning the sample in the described way went as follows:

- Kept schools in the 50 states plus DC (544 institutions kept)
- Dropped schools that didn't exist in the entirety of the time range, academic years 2003-2015 (505 institutions kept)
- Dropped schools that were missing bachelor's degrees awarded or graduation rate in any of the years in the data (457 institutions kept)
- Dropped schools that gave out 0 bachelor's degrees in a year or awarded fewer bachelor's degrees than associate degrees in a year (445 schools kept)
- Dropped schools that were missing values in all years from 2003 to 2015 for a few key variables (432 schools kept)

Dollar values were scaled to 2015 dollars using CPI, to allow for fair comparison, and monetary totals were scaled by total full time equivalent enrollment to account for the size of the institutions. There are alternate inflation indices, that may have been appropriate, for example HEPI - higher education price inflation, which is specific to the higher education sector, however these were not considered.

We also added state and year specific GDP and unemployment data in line with the approach in . The state GDP data came from the St. Louis Federal Reserve Bank's website for economic research (<https://fred.stlouisfed.org/>) and the state unemployment data came from the U.S. Department of Labor, Bureau of Labor Statistics website (<https://www.bls.gov/>).

Training and Testing Split

To validate our models we split our data into a test set and a training set. The test set was utilised only after models had been generated on the training set and did not inform our modeling decisions (e.g., such as dropping of variables)

Since each institution appears in our dataset 13 times (once for each year in 2003-2015), we clustered sampled by institution in splitting our dataset, selecting all years of data into the test or training set for each institution. In addition, we stratified by census division, to assure that each region would be represented in our training and testing set and reduce the variance in our estimators, as Tandberg and Hillman (2014) identified location as an important predictor of student outcomes. The census divisions were defined as below:

Table X: States in each Census Division

Region	States
New England	CT, ME, MA, NH, RI, VT
Middle Atlantic	NJ, NY, PA
East North Central	IN, IL, MI, OH, WI
West North Central	IA, KS, MN, MO, NE, ND, SD
South Atlantic	DE, DC, FL, GA, MD, NC, SC, VA, WV
East South Central	AL, KY, MS, TN
West South Central	AR, LA, OK, TX
Mountain	AZ, CO, ID, NM, MT, UT, NV, WY
Pacific	AK, CA, HI, OR, WA

We decided against including state fixed effects in any of our models, as we worried state fixed effects may account for too much variation, leading to poor prediction. We also decided against stratifying by state when splitting our data into a training and test set, and there were some states with only one institution in them, meaning states would be unrepresented in the training or test set.

After splitting the data, our training sample and test set looked as follows:

Table X: Training and Testing Split

	# of Institutions	# of Obs.
Training	347	4511
Testing	85	1105
Total	432	5616

Models Including Variables per Literature

Set-up

In attempt to hone in on the effect that changes in institution spending strategies have on student outcomes, we focused on the nine types of expenditure variables that were available to us in all institutions and years in our sample. These expenditure types are:

Table X: Expenditure Types and Examples

Expenditure Type	Example(s)
Instruction	Teacher Salaries
Public Service	Conferences, Community Services
Academic Support	Libraries, Museums, Demonstration Schools
Student Services	Student Activities/Organizations
Institutional Support	General Administration
Operation and Maintenance	Utilities
Depreciation	Cost of Captial Assets
Scholarships and Fellowships	Scholarships, Fellowships
Auxiliary Enterprises	Residence Halls, Student Health Services, Intercollegiate Sports

Along with these expenditure variables, which were converted to 2015 values and divided by FTE enrollment, we also included several institution and state characteristics as per the two research studies mentioned before. These extra covariates were:

Table X: Covariates Included as per Literature

<p style="text-align: center;"><u>State-level Variables</u></p> <p style="text-align: center;">GDP per Capita (state and year-specific) Unemployment Rate (state and year-specific)</p>
<p style="text-align: center;"><u>Institutional Characteristics</u></p> <p style="text-align: center;">Undergraduate Enrollment % of Total Enrollment that is Under-represented Minority (Black or Hispanic) Revenue from Pell Grants per FTE Revenue from State Appropriations per FTE Net Tuition and Fees Revenue per FTE Total Revenue per FTE</p>
<p style="text-align: center;"><u>Categorical Variables</u></p> <p style="text-align: center;">Year Dummies (2003-2015) Census Division Dummies (9 Divisions)</p>

We performed minor missing imputation on the expenditure variables and above covariates by imputing with the predicted values of a linear regression model including only dummies for institution and year.

Models

With these self-selected variables, we tried seven types of models on each of the two outcome variables:

- Traditional Linear Regression (LM)
- Traditional Linear Regression with Interactions
- Lasso Regression
- Lasso Regression with Interactions
- Ridge Regression
- Ridge Regression with Interactions
- Gaussian Kernel Regression

For the models with interactions, we interacted the nine expenditure variables with each other, as we hypothesized that spending more money in one expenditure area may change the effect of spending more money in another on these outcomes (e.g., spending more money on instruction may have a stronger effect if the institution has already invested heavily in making students feel comfortable on campus, such as high spending for residence halls or student health services).

For the Lasso and Ridge models, we performed ten-fold cross-validation on the training set to estimate the regularization coefficient (λ). For the Gaussian Kernel model, we chose the λ that minimized the leave-one-out cross-validation error on the whole training set.

The testing root mean squared error (RMSE) of the above models on both outcomes is provided in Table X below

Table X: Testing Raw RMSE of Self-Selected Models (standardized RMSE in parentheses)

Model	6-Year Graduation Rate	Bachelor's Degrees per FTE
LM	0.1008 (0.6385)	0.0348 (0.9194)
LM (with interactions)	0.0983 (0.6230)	0.0359 (0.9461)
Lasso	0.1004 (0.6359)	0.0346 (0.9129)
Lasso (with interactions)	0.0979 (0.6205)	0.0337 (0.8895)
Ridge	0.1005 (0.6366)	0.0344 (0.9074)
Ridge (with interactions)	0.0990 (0.6276)	0.0331 (0.8734)
Gaussian Kernel	0.0959 (0.6076)	0.0274 (0.7235)

For 6-year graduation rate, the gaussian kernel model had the lowest RMSE. However, the kernel model performed only slightly better than the Lasso model with interaction (the model with the second lowest testing RMSE), and in general, the linear models performed quite well on the test data compared to the kernel model. Among the linear models, the models with the interaction terms also performed better than the models without interactions, lending support to our hypothesis that there is some interaction effect between the expenditure variables.

For bachelor's degrees per FTE, the gaussian kernel model again had the lowest RMSE, and in this case performed much better than the linear models. And again, the linear models with interaction effects generally performed better than those without them. However, looking at the standardized errors in the above table (in parentheses), the models for bachelor's degrees per FTE performed much worse at predicting the outcome than those for 6-year graduation rate.

Interpretation

Because the kernel models performed the best on the testing set in terms of RMSE, we decided to interpret those models. In Table X below are the expenditure average partial derivatives from the kernel model that were significant at the 0.05 level for the 6-year graduation rate model.

Table X: Significant (at the 0.05 level) Average Partial Derivatives of Expenditure Variables for the 6-Year Graduation Rate Kernel Model

Expenditure	Estimate	Standard Error	p-value
Academic Support Expenditures per FTE	4.50e-06	1.20e-06	0.0001495
Student Services Expenditures per FTE	1.62e-05	1.70e-06	0
Institutional Support Expenditures per FTE	-1.25e-05	1.20e-06	0
Depreciation Expenditures per FTE	2.90e-06	7.00e-07	0.0000419
Auxiliary Enterprises Expenditures per FTE	2.07e-05	6.00e-07	0

For 6-year graduation rates, student services and auxiliary enterprises expenditures had the most positive average partial derivatives and instruction expenditures did not end up being significant. While we are careful to assign any causality from these models, given that all of the expenditure variables are included in these models as well as is the total dollars per FTE that the school has at its disposal in the year (Total Revenue per FTE), this is an interesting finding. It would be irresponsible to suggest that money spent to improve instruction does not have an effect on graduation rates (let alone the quality of education for all those at the school), and that all extra dollars should be funneled into student activities, residence halls, intercollegiate

sports, or health services. However, these models do suggest that efforts to improve students' experience and make students feel comfortable and supported on campus should be considered.

Automatic Variable Selection

Setup

For this approach the idea was to leave as much information available to the model as possible. Variables that were obviously colinear and not of interest were excluded, e.g. full time retention rate is a strong predictor of graduation rate. However this is not of interest, since it seems obvious that if students remain in full time enrollment, they will graduate within a specified time.

With of around 4000 observations, variables with more than 200 missing observations were dropped, with the remaining missing values imputed with a linear regression model including only dummies for region and year.

Selection

Since Lasso induces sparsity in the coefficient estimates, i.e it sets many of them to be 0, we used it to automatically select the variable with high predictive power for both our outcomes. We used 10 fold cross validation to select the optimal level for the regularisation parameters, 10 here is arbitrary though does not seem unreasonable since it is small in relation to our ~4000 observations.

We selected the covariates with the 10 variables with the largest absolute value coefficients, 10 here is arbitrary, though we suspect our models are robust to this since the size of the coefficients does decrease rapidly, due to the sparsity inducing Lasso.

Table X and X give the selected variables for each outcome as well as the value of their Lasso coefficients.

Lasso selected variables for graduation rate

Variable	Coefficient	Description
other_full_time	0.3876	Number of full-time undergrads not full-time, first-year
fed_grant_pct	-0.2412	Percentage of first, full-time undergrads with federal grants
total_part_time	-0.2004	Total number of part-time students
ptug_share_of_total_pt_enrl	-0.1639	Share of total part-time that are undergraduates
grant01	-0.1619	Total Pell Grants
auxiliary03	0.1453	Revenue from auxiliary enterprises
total_part_time_postbacc	-0.1141	Total number of part-time postbaccalaureate students
eandg02	0.0898	Total education and general expenditures - salaries and wages
bach_deg_share_of_tot_deg	0.08154	Share of total degrees at bachelor's level
fall_cohort_pct	0.08059	Fall cohort as percentage of total undergraduates

Lasso selected variables for bachelors per FTE

Variable	Coefficient	Description
bach_deg_share_of_tot_deg	0.5253	Share of total degrees at bachelor's level
hbcu2	0.4783	Indicator for historically black college or university status

Variable	Coefficient	Description
returning_to_total_undergraduate	0.3518	Share of total undergraduates completed first year
eandr_degree	-0.2552	Education and related expenses per degree
fed_grant_num	-0.159	Number of first, full-time undergrads with federal grants
other_ed_related_cost	0.1345	Instruction, student services, and other educational costs
total_enrollment_black_tot	-0.09466	Total enrollment of black students
inst_grant_num	0.09308	Number of first, full-time undergrads with institutional grants
ptug_share_of_total_pt_enrl	-0.08842	Share of total part-time that are undergraduates
credhoursug	0.08716	Instruction hours based on credit for undergraduates

The full definition of variables is given in the data dictionary publically available at <https://www.deltacostproject.org/>.

There are three main types of variable that the Lasso model has identified as strong predictors for both graduation rate and numbers of bachelors degrees per fte:

- part time variables
- expenditure variables
- grant variables

Part time variables are selected multiple times in both models, suggesting that there is an association between schools with higher portion of part time students and poorer student outcomes. Graduation rates only account for full time students. So we hypothesise that this link is identified due to institutions that are amenable to students dropping to part time study having a higher proportion of part timers. Such institutions have lower graduation rates since graduation rates only account for full time students. However interestingly it is also selected for the bachelors per FTE, which does account for part time students. This may suggest that part time students are less likely to graduate on time even when accounting for their reduced hours.

The selection of the expenditure variables, is promising and supports the hypothesis that the expenditures of institution are related to student outcomes.

We included grants in our first models due to the inclusion in Tandberg and Hillman (2014) and Slaughter et al. (2017), which use the grant variables as a proxy for deprivation of the student body, i.e. students from a lower socio - economic group are more likely to be awarded a federal grant. The inclusion of these variables for both outcomes in the automatic selection, suggests they are good predictors of outcomes, and suggests that including them in our first models did serve to reduce the variance of our estimators.

The signs of the coefficients for graduation rate are as expected, with high levels of part time study and federal grants having a negative impact and spending having a positive impact.

The coefficients for the Bachelor degrees per FTE are somewhat less interpretable, however the majority of the coefficients do follow the same signs as for graduation rates. However there is the notable exception of eandr_degree - educational spending per degree. We would expect this to have a positive impact not a negative impact. We think this is possibly due to high levels of correlation between one of the the selected variables and education spending resulting in the negative coefficient here to counteract this, this is likely due to the misspecification problems of the linear model. See below for our kernelized approach which helps to mitigate this problem.

Models

We used a classical, ridge, lasso and elastic net on the full data with all variables, and carried out kernel regularised regression with a Gaussian kernel on the variables selected by Lasso, for both outcomes. The results are summarised in the below tables:

Test set model performance for graduation rate

	Model	Standardized_RMSE
2	Ridge	0.2559
3	Lasso	0.2037
4	Elastic Net	0.2153
5	LM with lassoed variables with interactions	0.5413
6	Gaussian Kernel with lassoed variables	0.1878

Test set model performance for bachelors per FTE

	Model	Standardized_RMSE
2	Ridge	0.9688
3	Lasso	0.4710
4	Elastic Net	0.8979
5	LM with lassoed variables with interactions	9.8140
6	Gaussian Kernel with lassoed variables	0.3994

As expected the root mean square error for each of the models is lower than that achieved using the variables suggested in the literature. Using all the variables the Lasso performed best out of the regularised regression models for both outcomes.

We also note that the kernelized regression model performs the best overall for both outcomes, perhaps suggesting a complex non linear relationship between the selected variables.

Adding the interaction terms for the Lasso selected variables, with the hope of obtaining non trivial first difference estimates, essentially just resulted in overfitting, we could have mitigated this by adding another Lasso or ridge penalty term.

As in the first method, we analysed the Kernel model's average partial derivatives, in an attempt to understand the marginal effects of changes in the variables.

In general the Bachelors per FTE was a harder variable to predict with all the models performing less well, as such our interpretation on this variable holds less weight. This is consistent with the findings using the self selected variables.

Interpretation

Again, because the kernel models performed the best on the testing set in terms of RMSE, we decided to interpret those models. In Table X below are the average partial derivatives from the kernel model that were significant at the 0.05 level for the 6-year graduation rate model.

Table X: Top 10 Average Partial Derivatives of Lassoed Variables for the 6-Year Graduation Rate Kernel Model

Variable	Estimate (Standardised)	Standard Error	p-value
other full time	0.3928	0.011	0.0000
fed grant pct	-0.2004	0.0094	0.0000
grant01	-0.2792	0.011	0.0000
eandg02	0.1762	0.0099	0.0000
auxillary03	0.1442	0.01	0.0000
ptug share of total pt enrl	-0.1278	0.0074	0.0000
bach deg share of tot deg	0.1220	0.0077	0.0000
fall cohort pct	0.1073	0.0081	0.0000
total part time postbacc	-0.0625	0.0107	0.0000
total part time	-0.0966	0.0099	0.0000

Here it is noticeable out of the selected variables that the part time coefficients seem to have smaller effect sizes, with the exception of “other full time” which has the largest average partial derivative. In this model a small change in the “other full time” variable seems to account for the effect of a change in the part time proportion of an institution, since these variables are highly correlated in the data.

Federal grant have high negative average partial derivatives - agreeing with the Lasso model and our intuition, and our use of federal grants in our self selected variables.

Interestingly as with when we selected all expenditure variables, auxillary revenue - which is highly correlated with auxillary expenditure, had an important effect. However using this model education and related expenditure on salaries has a larger average partial derivative. We do not over interpret the fact that salaries were selected by the Lasso since this is highly correlated with the total expenditure variable.

Whereas our first approach identified this variable as having the largest average partial derivative. Thus there is some evidence that the tempting interpretation that auxillary expenditure is most important may be a model dependent and a step to far. However we do note that the fact that it was identified in both approaches as important does suggest that it has an effect. This in itself is interesting as it may be counterintuitive that spending not directly on education has an effect on the effectiveness of the institution as a whole, and perhaps justifies public spending on such activities, that could be deemed non essential when allocating public funds.

Table X: Top 10 Average Partial Derivatives of Lassoed Variables for the Bachelors per FTE Kernel Model

Variable	Estimate (Standardised)	Standard Error	p-value
hbcu2	0.2078	0.0400	0.0000
other ed related cost	0.0971	0.0027	0.0000
fed grant num	0.0836	0.0029	0.0000
inst grant num	0.0353	0.0023	0.0000
credhoursug	0.0291	0.0028	0.0000
eandr degree	-0.0282	0.0034	0.0000
returning to total undergraduate	0.0191	0.0025	0.0000
ptug share of total pt enrl	0.0106	0.0022	0.0000
total enrollment black tot	-0.0029	0.0036	0.4119

The average partial derivatives for the variables for the Bachelor’s degrees per FTE, aligns less well with our intuition and the first approach.

However we do note that part time effects seem to have diminished, having low average partial derivatives, this is somewhat expected since the Bachelor's per FTE outcome takes into account the level of part time studies at an institution. However, looking at the Lasso coefficients we might have inferred an association with part timers and poorer outcomes, despite the outcome accounting for part time study, however this seems not to be the case in light of the low coefficients in the kernel model for the part time variables.

Further interpretation is not obvious, though we note that the negative average partial derivative for eand r degree is still present, which may suggest that the reason for the negative lasso coefficient was not due to model mis specification, since the kernel model is much more flexible. Overall note that the effect size is much smaller here than it was for graduation rate, which matches with the Bachelor degrees per FTE being a noisier outcome variable.

Issues and Future Extensions

- All institutions treated equally i.e. no weighting based on size
- High levels of multicollinearity in expenditure variables, lasso may not be much better than random selection for variable selection
-

Conclusions

We make conclusions only in terms of interesting associations, and suggest more detailed research into areas of expenditure and student outcomes. However our models may suggest the following

- Auxilliary spending may worthwhile, since it is associated with positive outcomes
- Bachelor degrees per FTE is more unpredictable than graduation rate, we hypothesise since part time students have less predictable college careers.
-

References

- “Delta Cost Project.” 2018. Accessed February 15. <https://www.deltacostproject.org/>.
- Slaughter, A., K. Callahan, D. Kim, K. Meehan, L. Wainstein, L. Liu, and X. Zhu. 2017. “Playing the Long Game: The Impact of Outcomes-Based Funding on Bachelor’s Degree Attainment.” *Research for Action*.
- “St. Louis Federal Reserve Bank.” 2018. Accessed February 14. <https://fred.stlouisfed.org/>.
- Tandberg, and Hillman. 2014. “State Higher Education Performance Funding: Data, Outcomes, and Policy Implications.” *Journal of Education Finance* 39 (3). University of Illinois Press: 222–43.
- “U.S. Department of Labor, Bureau of Labor Statistics.” 2018. Accessed February 14. <https://www.bls.gov/>.